

Chapter 4

Apriori

Hiroshi Motoda and Kouzou Ohara

Contents

4.1	Introduction	62
4.2	Algorithm Description	62
4.2.1	Mining Frequent Patterns and Association Rules	62
4.2.1.1	Apriori	63
4.2.1.2	AprioriTid	66
4.2.2	Mining Sequential Patterns	67
4.2.2.1	AprioriAll	68
4.2.3	Discussion	69
4.3	Discussion on Available Software Implementations	70
4.4	Two Illustrative Examples	71
4.4.1	Working Examples	71
4.4.1.1	Frequent Itemset and Association Rule Mining	71
4.4.1.2	Sequential Pattern Mining	75
4.4.2	Performance Evaluation	76
4.5	Advanced Topics	80
4.5.1	Improvement in Apriori-Type Frequent Pattern Mining	80
4.5.2	Frequent Pattern Mining Without Candidate Generation	81
4.5.3	Incremental Approach	82
4.5.4	Condensed Representation: Closed Patterns and Maximal Patterns	82
4.5.5	Quantitative Association Rules	84
4.5.6	Using Other Measure of Importance/Interestingness	84
4.5.7	Class Association Rules	85
4.5.8	Using Richer Expression: Sequences, Trees, and Graphs	86
4.6	Summary	87
4.7	Exercises	88
	References	89

4.1 Introduction

Many of the pattern finding algorithms such as those for decision tree building, classification rule induction, and data clustering that are frequently used in data mining have been developed in the machine learning research community. Frequent pattern and association rule mining is one of the few exceptions to this tradition. Its introduction boosted data mining research and its impact is tremendous. The basic algorithms are simple and easy to implement. In this chapter the most fundamental algorithms of frequent pattern and association rule mining, known as Apriori and AprioriTid [3, 4], and Apriori's extension to sequential pattern mining, known as AprioriAll [6, 5], are explained based on the original papers with working examples, and performance analysis of Apriori is shown using a freely available implementation [1] for a dataset in UCI repository [8]. Since Apriori is so fundamental and the form of database is limited to market transaction, there have been many works for improving computational efficiency, finding more compact representation, and extending the types of data that can be handled. Some of the important works are also briefly described as advanced topics.

4.2 Algorithm Description

4.2.1 Mining Frequent Patterns and Association Rules

One of the most popular data mining approaches is to find frequent itemsets from a transaction dataset and derive association rules. The problem is formally stated as follows. Let $\mathcal{I} = \{i_1, i_2, \dots, i_m\}$ be a set of items. Let \mathcal{D} be a set of transactions, where each transaction t is a set of items such that $t \subseteq \mathcal{I}$. Each transaction has a unique identifier, called its *TID*. A transaction t contains X , a set of some items in \mathcal{I} , if $X \subseteq t$. An association rule is an implication of the form $X \Rightarrow Y$, where $X \subset \mathcal{I}$, $Y \subset \mathcal{I}$, and $X \cap Y = \emptyset$. The rule $X \Rightarrow Y$ holds in \mathcal{D} with confidence c ($0 \leq c \leq 1$) if the fraction of transactions that also contain Y in those which contain X in \mathcal{D} is c . The rule $X \Rightarrow Y$ (and equivalently $X \cup Y$) has support¹ s ($0 \leq s \leq 1$) in \mathcal{D} if the fraction of transactions in \mathcal{D} that contain $X \cup Y$ is s . Given a set of transactions \mathcal{D} , the problem of mining association rules is to generate all association rules that have support and confidence no less than the user-specified minimum support (called *minsup*) and minimum confidence (called *minconf*), respectively.

Finding frequent² itemsets (itemsets with support no less than *minsup*) is not trivial because of the computational complexity due to combinatorial explosion. Once

¹An alternative support definition is the absolute count of frequency. In this chapter the latter definition is also used where appropriate.

²The Apriori paper [3] uses "large" to mean "frequent," but large is often associated with the number of items in the itemset. Thus, we prefer to use "frequent."

frequent itemsets are obtained, it is straightforward to generate association rules with confidence no less than *minconf*. Apriori and AprioriTid, proposed by R. Agrawal and R. Srikant, are seminal algorithms that are designed to work for a large transaction dataset [3].

4.2.1.1 Apriori

Apriori is an algorithm to find all sets of items (itemsets) that have support no less than *minsup*. The support for an itemset is the ratio of the number of transactions that contain the itemset to the total number of transactions. Itemsets that satisfy minimum support constraint are called *frequent itemsets*. Apriori is characterized as a level-wise complete search (breadth first search) algorithm using anti-monotonicity property of itemsets: “If an itemset is not frequent, any of its superset is never frequent,” which is also called the *downward closure property*. The algorithm makes multiple passes over the data. In the first pass, the support of individual items is counted and frequent items are determined. In each subsequent pass, a seed set of itemsets found to be frequent in the previous pass is used for generating new potentially frequent itemsets, called *candidate itemsets*, and their actual support is counted during the pass over the data. At the end of the pass, those satisfying minimum support constraint are collected, that is, frequent itemsets are determined, and they become the seed for the next pass. This process is repeated until no new frequent itemsets are found.

By convention, Apriori assumes that items within a transaction or itemset are sorted in lexicographic order. The number of items in an itemset is called its *size* and an itemset of size k is called a k -itemset. Let the set of frequent itemsets of size k be F_k and their candidates be C_k . Both F_k and C_k maintain a field, support count.

Apriori algorithm is given in Algorithm 4.1. The first pass simply counts item occurrences to determine the frequent 1-itemsets. A subsequent pass consists of two phases. First, the frequent itemsets F_{k-1} found in the $(k - 1)$ -th pass are used to generate the candidate itemsets C_k using the apriori-gen function. Next, the database is scanned and the support of candidates in C_k is counted. The subset function is used for this counting.

The apriori-gen function takes as argument F_{k-1} , the set of all frequent $(k - 1)$ -itemsets, and returns a superset of the set of all frequent k -itemsets. First, in the join steps, F_{k-1} is joined with F_{k-1} .

insert into C_k

select $p.\text{fitemset}_1, p.\text{fitemset}_2, \dots, p.\text{fitemset}_{k-1}, q.\text{fitemset}_{k-1}$

from $F_{k-1}p, F_{k-1}q$

where $p.\text{fitemset}_1 = q.\text{fitemset}_1, \dots, p.\text{fitemset}_{k-2} = q.\text{fitemset}_{k-2},$
 $p.\text{fitemset}_{k-1} < q.\text{fitemset}_{k-1}$

Here, $F_k p$ means that the itemset p is a frequent k -itemset, and $p.\text{fitemset}_k$ is the k -th item of the frequent itemset p .

Then, in the prune step, all the itemsets $c \in C_k$ for which some $(k - 1)$ -subset is not in F_{k-1} are deleted.

Algorithm 4.1 Apriori Algorithm

```

 $F_1 = \{\text{frequent 1-itemsets}\};$ 
for ( $k = 2; F_{k-1} \neq \emptyset; k++$ ) do begin
   $C_k = \text{apriori-gen}(F_{k-1});$  //New candidates
  foreach transaction  $t \in \mathcal{D}$  do begin
     $C_t = \text{subset}(C_k, t);$  //Candidates contained in  $t$ 
    foreach candidate  $c \in C_t$  do
       $c.\text{count}++;$ 
    end
     $F_k = \{c \in C_k \mid c.\text{count} \geq \text{minsup}\};$ 
  end
  Answer =  $\cup_k F_k;$ 

```

The subset function takes as arguments C_k and a transaction t , and returns all the candidate itemsets contained in the transaction t . For fast counting, Apriori adopts a hash-tree to store the candidate itemsets C_k . Itemsets are stored in leaves. Every node is initially a leaf node, and the depth of the root node is defined to be 1. When the number of itemsets in a leaf node exceeds a specified threshold, the leaf node is converted to an interior node. An interior node at depth d points to nodes at depth $d + 1$. Which branch to follow is decided by applying a hash function to the d -th item of the itemset. Thus, each leaf node is ensured to contain at most a certain number of itemsets (to be precise, this is true only when creating an interior node takes place at depth d smaller than k), and an itemset in the leaf node can be reached by successively hashing each item in the itemset in sequence from the root. Once the hash-tree is constructed, the subset function finds all the candidates contained in a transaction t , starting from the root node. At the root node, every item in t is hashed, and each branch determined is followed one depth down. If a leaf node is reached, itemsets in the leaf that are in the transaction t are searched and those found are made reference to the answer set. If an interior node is reached by hashing the item i , items that come after i in t are hashed recursively until a leaf node is reached. It is evident that itemsets in the leaves that are never reached are not contained in t .

Clearly, any subset of a frequent itemset satisfies the minimum support constraint. The join operation is equivalent to extending F_{k-1} with each item in the database and then deleting those itemsets for which the $(k - 1)$ -itemset obtained by deleting the $(k - 1)$ -th item is not in F_{k-1} . The condition $p.\text{itemset}_{k-1} < q.\text{itemset}_{k-1}$ ensures that no duplication is made. The prune step where all the itemsets whose $(k - 1)$ -subsets are not in F_{k-1} are deleted from C_k does not delete any itemset that could be in F_k . Thus, $C_k \supseteq F_k$, and Apriori algorithm is correct.

The remaining task is to generate the desired association rules from the frequent itemsets. A straightforward algorithm for this task is as follows. To generate rules,

all nonempty subsets of every frequent itemset f are enumerated and for every such subset a , a rule of the form $a \Rightarrow (f - a)$ is generated if the ratio of $\text{support}(f)$ to $\text{support}(a)$ is at least minconf . Here, note that the confidence of the rule $\hat{a} \Rightarrow (f - \hat{a})$ cannot be larger than the confidence of $a \Rightarrow (f - a)$ for any $\hat{a} \subset a$. This in turn means that for a rule $(f - a) \Rightarrow a$ to hold, all rules of the form $(f - \hat{a}) \Rightarrow \hat{a}$ must hold. Using this property, the algorithm to generate association rules is given in Algorithm 4.2.

Algorithm 4.2 Association Rule Generation Algorithm

```

 $H_1 = \emptyset$  //Initialize
foreach; frequent  $k$ -itemset  $f_k, k \geq 2$  do begin
   $A = (k - 1)$ -itemsets  $a_{k-1}$  such that  $a_{k-1} \subset f_k$ ;
  foreach  $a_{k-1} \in A$  do begin
     $\text{conf} = \text{support}(f_k) / \text{support}(a_{k-1})$ ;
    if ( $\text{conf} \geq \text{minconf}$ ) then begin
      output the rule  $a_{k-1} \Rightarrow (f_k - a_{k-1})$ 
        with confidence =  $\text{conf}$  and support =  $\text{support}(f_k)$ ;
      add  $(f_k - a_{k-1})$  to  $H_1$ ;
    end
  end
  call ap-genrules( $f_k, H_1$ );
end

Procedure ap-genrules( $f_k$ : frequent  $k$ -itemset,  $H_m$ : set of  $m$ -item
  consequents)
if ( $k > m + 1$ ) then begin
   $H_{m+1} = \text{apriori-gen}(H_m)$ ;
  foreach  $h_{m+1} \in H_{m+1}$  do begin
     $\text{conf} = \text{support}(f_k) / \text{support}(f_k - h_{m+1})$ ;
    if ( $\text{conf} \geq \text{minconf}$ ) then
      output the rule  $f_k - h_{m+1} \Rightarrow h_{m+1}$ 
        with confidence =  $\text{conf}$  and support =  $\text{support}(f_k)$ ;
    else
      delete  $h_{m+1}$  from  $H_{m+1}$ ;
    end
  call ap-genrules( $f_k, H_{m+1}$ );
end

```

Apriori achieves good performance by reducing the size of candidate sets. However, in situations with very many frequent itemsets or very low minimum support, it still suffers from the cost of generating a huge number of candidate sets and scanning the database repeatedly to check a large set of candidate itemsets.

4.2.1.2 AprioriTid

AprioriTid is a variation of Apriori. It does not reduce the number of candidates but it does not use the database \mathcal{D} for counting support after the first pass. It uses a new dataset $\overline{\mathcal{C}}_k$. Each member of the set $\overline{\mathcal{C}}_k$ is of the form $\langle TID, \{ID\} \rangle$, where each ID is the identifier of a potentially frequent k -itemset present in the transaction with identifier TID except $k = 1$. For $k = 1$, $\overline{\mathcal{C}}_1$ corresponds to the database \mathcal{D} , although conceptually each item i is replaced by the itemset $\{i\}$. The member of $\overline{\mathcal{C}}_k$ corresponding to a transaction t is $\langle t.TID, \{c \in C_k | c \text{ contained in } t\} \rangle$.

The intuition for using $\overline{\mathcal{C}}_k$ is that it will be smaller than the database \mathcal{D} for large values of k because some transactions may not contain any candidate k -itemset, in which case $\overline{\mathcal{C}}_k$ does not have an entry for this transaction, or because very few candidates may be contained in the transaction and each entry may be smaller than the number of items in the corresponding transaction. AprioriTid algorithm is given in Algorithm 4.3. Here, $c[i]$ represents the i -th item in k -itemset c .

Algorithm 4.3 AprioriTid Algorithm

```

 $F_1 = \{\text{frequent 1-itemsets}\};$ 
 $\overline{\mathcal{C}}_1 = \text{database } \mathcal{D};$ 
for ( $k = 2; F_{k-1} \neq \emptyset; k++$ ) do begin
     $C_k = \text{apriori-gen}(F_{k-1});$  //New candidates
     $\overline{\mathcal{C}}_k = \emptyset;$ 
    foreach entry  $t \in \overline{\mathcal{C}}_{k-1}$  do begin
        // determine candidate itemsets in  $C_k$  contained
        // in the transaction with identifier  $t.TID$ 
         $C_t = \{c \in C_k | (c - c[k]) \in t.\text{set-of-itemsets} \wedge$ 
             $(c - c[k - 1]) \in t.\text{set-of-itemsets}\};$ 
        foreach candidate  $c \in C_t$  do
             $c.\text{count}++;$ 
            if ( $C_t \neq \emptyset$ ) then  $\overline{\mathcal{C}}_k += \langle t.TID, C_t \rangle;$ 
    end
     $F_k = \{c \in C_k \mid c.\text{count} \geq \text{minsup}\};$ 
end
Answer =  $\cup_k F_k;$ 

```

Each $\overline{\mathcal{C}}_k$ is stored in a sequential structure. A candidate k -itemset c_k in C_k maintains two additional fields; generator and extensions, in addition to the field, support count. The generator field stores the IDs of the two frequent $(k - 1)$ -itemsets whose join generated c_k . The extension field stores the IDs of all the $(k + 1)$ -candidates that are extensions of c_k . When a candidate c_k is generated by joining f_{k-1}^1 and f_{k-1}^2 , their IDs are saved in the generator field of c_k and the ID of c_k is added to the extension field of f_{k-1}^1 . The $t.\text{set-of-itemsets}$ field of an entry t in $\overline{\mathcal{C}}_{k-1}$ gives the IDs of all

$(k - 1)$ -candidates contained in $t.TID$. For each such candidate c_{k-1} the extension field gives T_k , the set of ID s of all the candidate k -itemsets that are extensions of c_{k-1} . For each c_k in T_k , the generator field gives the ID s of the two itemsets that generated c_k . If these itemsets are present in the entry for $t.set$ -of-itemsets, it is concluded that c_k is present in transaction $t.TID$, and c_k is added to C_t .

AprioriTid has an overhead to calculate \overline{C}_k but an advantage that \overline{C}_k can be stored in memory when k is large. It is thus expected that Apriori beats AprioriTid in earlier passes (small k) and AprioriTid beats Apriori in later passes (large k). Since both Apriori and AprioriTid use the same candidate generation procedure and therefore count the same itemsets, it is possible to make a combined use of these two algorithms in sequence. AprioriHybrid uses Apriori in the initial passes and switches to AprioriTid when it expects that the set \overline{C}_k at the end of the pass will fit in memory.

4.2.2 Mining Sequential Patterns

Agrawal and Srikant extended Apriori algorithm to the problem of sequential pattern mining [6]. In Apriori there is no notion of sequence, and thus, the problem of finding which items appear together can be viewed as finding intratransaction patterns. Here, sequence matters and the problem of finding sequential patterns can be viewed as intertransaction patterns.

Each transaction consists of sequence-id, transaction-time, and a set of items. The same sequence-id has no more than one transaction with the same transaction-time. A *sequence* is an ordered list of itemsets. Thus, a sequence consists of a list of sets of characters (items), rather than being simply a list of characters. The length of a sequence is the number of itemsets in the sequence. A sequence of length k is called a k -sequence. Without loss of generality, the set of items is assumed to be mapped to a set of contiguous integers, and an itemset i is denoted by $(i_1 i_2 \dots i_m)$ where i_j is an item. A sequence s is denoted by $\langle s_1 s_2 \dots s_n \rangle$. A sequence $\langle a_1 a_2 \dots a_n \rangle$ is contained in another sequence $\langle b_1 b_2 \dots b_m \rangle$ ($n \leq m$) if there exist integers $i_1 < i_2 < \dots < i_n$ such that $a_1 \subseteq b_{i_1}, a_2 \subseteq b_{i_2}, \dots, a_n \subseteq b_{i_n}$. All the transactions with the same sequence-id which are sorted by transaction-time together form a sequence (transaction sequence). A sequence-id supports a sequence s if s is contained in its transaction sequence. The support for a sequence is defined as the fraction of total number of sequence-ids that support this sequence. Likewise, the support for an itemset i is defined as the fraction of sequence-ids that have items in i in any one of their transactions. Note that this definition is different from that used in Apriori. Thus the itemset i and the 1-sequence $\langle i \rangle$ have the same support.

Given a transaction database \mathcal{D} , the problem of mining sequential patterns is to find the maximal³ sequences among all sequences that satisfy a certain user-specified minimum support constraint. Each such maximal sequence represents a sequential pattern. A sequence satisfying the minimum support constraint is called a *frequent sequence* (not necessarily maximal), and an itemset satisfying the minimum support

³Later R. Agrawal and R. Srikant removed this constraint in their generalized sequential patterns (GSP) [32].

constraint is called a frequent itemset, or fitemset for short. Any frequent sequence must be a list of fitemsets.

The algorithm consists of five phases: (1) sort phase, (2) fitemset phase, (3) transformation phase, (4) sequence phase, and (5) maximal phase. The first three are preprocessing phases and the last one is a postprocessing phase.

In the sort phase, the database \mathcal{D} is sorted with sequence-id as the major key and transaction-time as the minor key. In the fitemset phase, the set of all fitemsets is obtained using Apriori algorithm with the corresponding modification of counting a support, and is mapped to a set of contiguous integers. This makes comparing two fitemsets for equality in a constant time. Note that the set of all frequent 1-sequences are simultaneously found in this phase. In the transformation phase, each transaction is replaced by the set of all fitemsets that are in that transaction. If a transaction does not contain any fitemset, it is not retained in the transformed sequence. If a transaction sequence does not contain any fitemset, this sequence is removed from the transformed database, but it is still used in counting the total number of sequence-ids. After the transformation, a transaction sequence is represented by a list of sets of fitemsets. Each set of fitemsets is represented by $\{f_1, f_2, \dots, f_n\}$, where f_i is an fitemset. This transformation is designed for efficiently testing which given frequent sequences are contained in a transaction sequence. The transformed database is denoted as \mathcal{D}_T .

The sequence phase is the main part where the frequent sequences are to be enumerated. Two families of algorithms are proposed: count-all and count-some. They differ in the way the frequent sequences are counted. Count-all algorithm counts all the frequent sequences, including nonmaximal sequences that must be pruned later, whereas count-some algorithm avoids counting sequences which are contained in a longer sequence because the final goal is to obtain only maximal sequences. Agrawal and Srikant developed one count-all algorithm called AprioriAll and two count-some algorithms called AprioriSome and DynamicSome. Here, only AprioriAll is explained due to the space limitation.

In the last maximal phase, maximal sequences are extracted from the set of all frequent sequences. The hash-tree (similar to the one used in the subset function in Apriori) is used to quickly find all subsequences of a given sequence.

4.2.2.1 AprioriAll

The algorithm is given in Algorithm 4.4. In each pass the frequent sequences from the previous pass are used to generate the candidate sequences and then their support is measured by making a pass over the database. At the end of the pass, the support of the candidates is used to determine the frequent sequences.

The apriori-gen-2 function takes as argument F_{k-1} , the set of all frequent $(k-1)$ -sequences. First, join operation is performed as

```
insert into  $C_k$ 
select  $p.\text{fitemset}_1, p.\text{fitemset}_2, \dots, p.\text{fitemset}_{k-1}, q.\text{fitemset}_{k-1}$ 
from  $F_{k-1}p, F_{k-1}q$ 
where  $p.\text{fitemset}_1 = q.\text{fitemset}_1, \dots, p.\text{fitemset}_{k-2} = q.\text{fitemset}_{k-2},$ 
```


Algorithm 4.4 AprioriAll Algorithm

```

 $F_1 = \{\text{frequent 1-sequences}\};$  // Result of fitemset phase
for ( $k = 2$ ;  $F_{k-1} \neq \emptyset$ ;  $k++$ ) do begin
   $C_k = \text{apriori-gen-2}(F_{k-1});$  //New candidate sequences
  foreach transaction sequence  $t \in \mathcal{D}_T$  do begin
     $C_t = \text{subseq}(C_k, t);$  //Candidate sequences contained in  $t$ 
    foreach candidate  $c \in C_t$  do
       $c.\text{count}++;$ 
    end
     $F_k = \{c \in C_k \mid c.\text{count} \geq \text{minsup}\};$ 
  end
  Answer = maximal sequences in  $\cup_k F_k$ ;

```

then, all the sequences $c \in C_k$ for which some $(k - 1)$ -subsequence is not in F_{k-1} are deleted. The subseq function is similar to the subset function in Apriori. As in Apriori, the candidate sequences C_k are stored in a hash-tree to quickly find all candidates contained in a transaction sequence. Note that the transformed transaction sequence is a list of sets of fitemsets and all the fitemsets in a set have the same transaction-time, and no more than one transaction with the same transaction-time is allowed for the same sequence-id. This constraint has to be imposed in the subseq function.

4.2.3 Discussion

Both Apriori and AprioriTid need *minsup* and *minconf* to be specified in advance. The algorithms have to be rerun each time these values are changed, throwing everything away that was obtained in previous runs. If no appropriate values for these thresholds are known in advance and we want to know how the results change with these values without rerunning the algorithms, the best we can do is to generate and count only those itemsets that appear at least once in the database without duplication and store them all in an efficient way. Note that Apriori generates candidates that do not exist in the database.

Apriori and AprioriTid use a hash-tree to store the candidate itemsets. Another data structure that is often used is a trie-structure [35, 9]. Each node in the depth k of the trie corresponds to a candidate k -itemset and stores the k -th item and the support of the itemset. As two frequent k -itemsets that share the first $(k - 1)$ -itemsets are siblings below their parent node at the depth $k - 1$ in the trie, the candidate generation is simply to join the two siblings, and extend the tree to one more depth below the first frequent k -itemset after pruning. In order to find the candidate k -itemsets that are contained in a transaction t , each item in the transaction is fed from the root node and the branch is followed according to the succeeding item until a k -th item is reached. Many practical implementations of Apriori use this trie-structure to store not only candidates but also transactions [10, 9].

If we go a step further, we can get rid of generating candidate itemsets at all. Further, it is not necessary to enumerate all the frequent itemsets. These topics are discussed in Section 4.5.

Apriori and almost all other association rule minings use two-phase strategy: first mine frequent patterns and then generate association rules. This is not the sole way. Webb's MagnumOpus uses another strategy that immediately generates a large subset of all association rules [38].

There are direct extensions of the original Apriori family. Use of taxonomy and incorporating temporal constraint are two examples. Generalized association rules [30] employ a set of user-specified taxonomies, which makes it possible to extract frequent itemsets that are expressed by higher concepts even when use of the base level concepts produces only infrequent itemsets. The basic algorithm is to add all ancestors of each item in a transaction to the transaction and then run Apriori algorithm. Several optimizations can be added to improve efficiency, one example being that the support for an itemset X that contains both an item x and its ancestor \hat{x} is the same as the support of the itemset $X - \hat{x}$, and thus need not be counted. Generalized sequential patterns [32] place, in addition to the introduction of taxonomies, time constraints that specify a minimum and/or maximum time period between adjacent elements (itemsets) in a pattern and relax the restrictions that items in an element of a sequential pattern must come from the same transaction by allowing the items to be present in a set of transactions of the same sequence-id whose transaction-times are within a user-specified time window. It also finds all frequent sequential patterns (not limited to maximal sequential patterns). GSP algorithm runs about 20 times faster than AprioriAll, one reason being that GSP counts fewer candidates than AprioriAll.

4.3 Discussion on Available Software Implementations

There are many available implementations of Apriori ranging from free software to commercial products. Here, we will present only three well-known implementations which are freely downloadable via Internet.

The first one is an implementation embedded in the most famous open-source machine learning and data mining toolkit, Weka, provided by the University of Waikato [40]. Apriori in Weka can be used through Weka's common graphical user interface together with many other algorithms that are available in Weka. The implementation includes Weka's own extensions. For example, *minsup* is iteratively decreased from an upper bound U_{minsup} to a lower bound L_{minsup} with an interval δ_{minsup} . Further, in addition to confidence the metrics lift, leverage, and conviction are available to evaluate association rules. Lift and leverage are discussed in Section 4.5. Conviction [11] is a metric that was proposed to measure the departure from independence of an association rule taking implication into account. When using one of these metrics, its minimal value has to be given as a threshold.

The second implementation is the one by Christian Borgelt [1], which is distributed under the terms of the GNU Lesser (Library) General Public License. This implementation is basically a command line application, and some graphical user interfaces are separately available. It essentially follows the flow of the original Apriori, but has its own extensions, too, to make it faster and to reduce its memory use. It employs a trie called the prefix tree to store both transactions and itemsets for efficient support counting [10]. The prefix tree is slightly different from the trie explained in Subsection 4.2.3. Optionally, the user can choose to use a simple list instead of a prefix tree to store transactions. Furthermore, this implementation can find not only frequent itemsets and association rules, but also closed itemsets, and maximal itemsets. Closed and maximal itemsets are discussed in Section 4.5. In addition, several metrics other than confidence, such as information gain, are also available in this implementation to evaluate and select association rules.

The third implementation is the one by Fence Bodon, which is freely distributed for research purposes [2]. This implementation is also trie-based, similar to Borgelt's, but adopts a trie with a simpler structure, and computes only frequent itemsets and association rules. It works as a command line application, and accepts four arguments. The first three are mandatory: an input file, including transactions, an output file, and *minsup*. The fourth is *minconf*, which is optional. If *minconf* is given, association rules are mined, as well as frequent itemsets; otherwise, it outputs only frequent itemsets. This implementation is written in C++ to provide object-oriented components which can be easily reused to develop other Apriori-based algorithms.

4.4 Two Illustrative Examples

4.4.1 Working Examples

We will illustrate the detailed behavior of the aforementioned algorithms using a small database shown in Table 4.1, where SID and TT mean the sequence-id and transaction-time, respectively. We use this database in both association rule (frequent itemset) mining and maximal sequential pattern mining. In the former case SID and TT are ignored.

4.4.1.1 Frequent Itemset and Association Rule Mining

Suppose that we want to find frequent itemsets under *minsup* = 0.2 and association rules with *minconf* = 0.6.

Apriori (Algorithm 4.1)

Apriori first scans the whole database and derives a set of frequent 1-itemsets appearing in at least three transactions, $F_1 = \{a, c, d, f, g\}$. From this F_1 , the apriori-gen function derives a set of candidate frequent 2-itemsets $C_2 = \{ac, ad, af, ag, cd, cf, cg, df, dg, fg\}$. C_2 consists of all possible pairs of elements of F_1 since no pruning is made at this stage.

TABLE 4.1 A Transaction Database of the Working Example

TID	SID	TT	Items
001	1	May 03	<i>c, d</i>
002	1	May 05	<i>f</i>
003	4	May 05	<i>a, c</i>
004	3	May 05	<i>c, d, f</i>
005	2	May 05	<i>b, c, f</i>
006	3	May 06	<i>d, f, g</i>
007	4	May 06	<i>a</i>
008	4	May 07	<i>a, c, d</i>
009	3	May 08	<i>c, d, f, g</i>
010	1	May 08	<i>d, e</i>
011	2	May 08	<i>b, d</i>
012	3	May 09	<i>d, g</i>
013	1	May 09	<i>e, f</i>
014	3	May 10	<i>c, d, f</i>

Next, Apriori computes their support by scanning the database using the subset function, which utilizes a hash-tree. Figure 4.1 briefly illustrates how a hash-tree is constructed and used. Suppose that the elements of C_2 are added into the hash-tree in lexicographic order, and the maximum number of itemsets allowed to be in a leaf node is 4. Thus, the number of itemsets in the root (leaf) node exceeds the threshold when the fifth itemset cd is to be added. Then, the node is converted into an interior

In case that the maximum number of itemsets that can be stored in a node is 4.

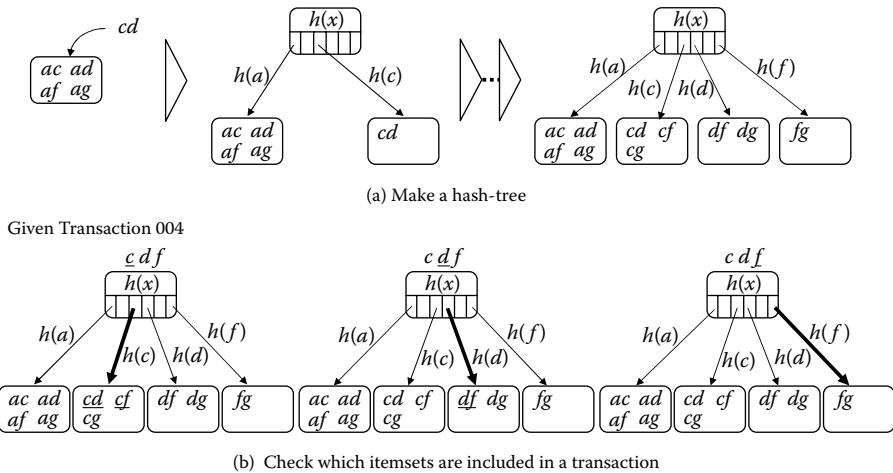


Figure 4.1 Example of hash-tree.

one, and each itemset branches into the corresponding new leaf node according to the hash value given by the function $h(x)$, where x is an item, the first item in each itemset in this case. We assume that $h(x)$ is given in advance and is common for all nodes. Since the first four itemsets share the same first item a , they fall into the same leaf node, while cd falls into a different one. When checking which of the candidates are included in a transaction, for example, Transaction 004, each item in the transaction is hashed at the root node. For example, by hashing c in cdf , it reaches the second left leaf node, and two itemsets cd and cf are found to be subsets of cdf as shown in the left tree of Figure 4.1(b). Next, by hashing d , df is found in the third left leaf node (the middle tree), but by hashing f , no subset of cdf is found in the rightmost leaf node (the right tree). As a result, the support counts of these itemsets found, cd , cf , and df , are increased by 1. Note that, after all the transactions have been processed, the frequencies of the candidates af and ag are found to be 0. This means that Apriori may generate candidates that do not exist in a given database.

After this support counting, $F_2 = \{cd, cf, df, dg\}$ is derived. These frequent 2-itemsets in F_2 are used as the seeds of frequent 3-itemsets. The itemsets cd and cf in F_2 sharing the first item c are joined and yield a new candidate cdf by apriori-gen because df is also included in F_2 . The itemsets df and dg are also joined as well, but the resulting candidate is pruned because its subset fg is not included in F_2 . Consequently, C_3 , a set of candidate frequent 3-itemsets, consists of cdf only. Then, Apriori counts its support by scanning the database again, and derives $F_3 = \{cdf\}$. No candidate frequent 4-itemsets can be generated from this F_3 because it contains only one itemset. Thus, Apriori terminates.

AprioriTid (Algorithm 4.3)

Apriori has to scan the whole database three times to obtain these frequent itemsets, but AprioriTid (Algorithm 4.3) scans it only once for the first pass, and makes and uses new datasets \overline{C}_1 and \overline{C}_2 to count the support of candidates in C_2 and C_3 , respectively. Figure 4.2 illustrates how AprioriTid finds frequent itemsets from these datasets. \overline{C}_2 is generated while counting the support of each candidate in C_2 , whereas \overline{C}_1 is generated directly from the given database. Suppose $t = \langle 001, \{\{c\}, \{d\}\} \rangle \in \overline{C}_1$. Then, a candidate cd in C_2 is added to C_t because t .set-of-itemsets ($\{\{c\}, \{d\}\}$) contains both 1-itemsets constituting cd . More precisely, cd is added to C_t because it is a union of two 1-itemsets in t , which means Transaction 001 supports cd . No other candidate is added to C_t as Transaction 001 does not support any other candidate in C_2 . Then, the support count of cd is increased by 1, and $\langle 001, \{\{cd\}\} \rangle$ is added to \overline{C}_2 . Similarly, $\langle 003, \{\{ac\}\} \rangle$ is added to \overline{C}_2 because Transaction 003 supports $ac \in C_2$, although an entry corresponding to $\langle 002, \{\{f\}\} \rangle$ of \overline{C}_1 is not because Transaction 002 does not support any 2-itemsets. Eventually, \overline{C}_2 has 9 entries, as shown in Figure 4.2, whose size is smaller than that of the given database. \overline{C}_3 is generated in the same manner during the support counting of candidates in C_3 . Since the unique candidate in C_3 is cdf , only the three entries of \overline{C}_2 , including both cd and cf , whose union is cdf , survive in \overline{C}_3 . Note that \overline{C}_3 is generated, but actually never used because C_4 becomes empty.




Figure 4.2 Example of AprioriTid.

Association rules (Algorithm 4.2)

Next, association rules are generated from the found frequent itemsets according to Algorithm 4.2 for the given $minconf = 0.6$. Let us consider frequent 2-itemsets, cd , cf , df , and dg , first. It is obvious that only two kinds of rules can be generated from each itemset. Table 4.2 summarizes the resulting rules and their confidence. The association rules 1 and 8 are the outputs by Algorithm 4.2 because they satisfy the $minconf$ constraint. The procedure $ap_genrules$ is called for each of these satisfactory rules, but it outputs nothing because it no longer generates other rules from the 2-itemsets.

Then, Algorithm 4.2 tries to generate association rules from the frequent 3-itemset, cdf . First, it generates three association rules with 1-item consequent as shown in the left half of Table 4.3. Algorithm 4.2 returns all of them as they satisfy the $minconf$ constraint. After that, the procedure $ap_genrules$ is called, taking cdf and $\{c, d, f\}$

TABLE 4.2 Association Rules Generated from Frequent 2-Itemsets

No.	Rule	Confidence	No.	Rule	Confidence
1	$c \Rightarrow d$	0.71	5	$d \Rightarrow f$	0.44
2	$d \Rightarrow c$	0.56	6	$f \Rightarrow d$	0.57
3	$c \Rightarrow f$	0.57	7	$d \Rightarrow g$	0.33
4	$f \Rightarrow c$	0.57	8	$g \Rightarrow d$	1.0

TABLE 4.3 Association Rules Generated from Frequent 3-Itemsets

1-Item Consequent			2-Item Consequent		
No.	Rule	Confidence	No.	Rule	Confidence
9	$cd \Rightarrow f$	0.60	12	$f \Rightarrow cd$	0.43
10	$cf \Rightarrow d$	0.75	13	$d \Rightarrow cf$	0.33
11	$df \Rightarrow c$	0.75	14	$c \Rightarrow df$	0.43

as its arguments. A set of 2-itemsets $\{cd, cf, df\}$ is derived by the function *apriori-gen* called within *ap-genrules*, each of which is used as the consequent of a new association rule. The resulting three rules are shown in the right half of Table 4.3. But, none of them can be the outputs because their confidence is less than the specified $minconf = 0.6$. Since 3-item consequents cannot be obtained from cdf , *ap-genrules* terminates, and Algorithm 4.2 terminates too because $F_4 = \emptyset$.

4.4.1.2 Sequential Pattern Mining

Next, we find frequent maximal sequential patterns from the same transaction database in Table 4.1 by using *AprioriAll* (Algorithm 4.4) for $minsup = 0.3$. Figure 4.3 illustrates the flow of the first three phases, that is, sort phase, itemset phase, and

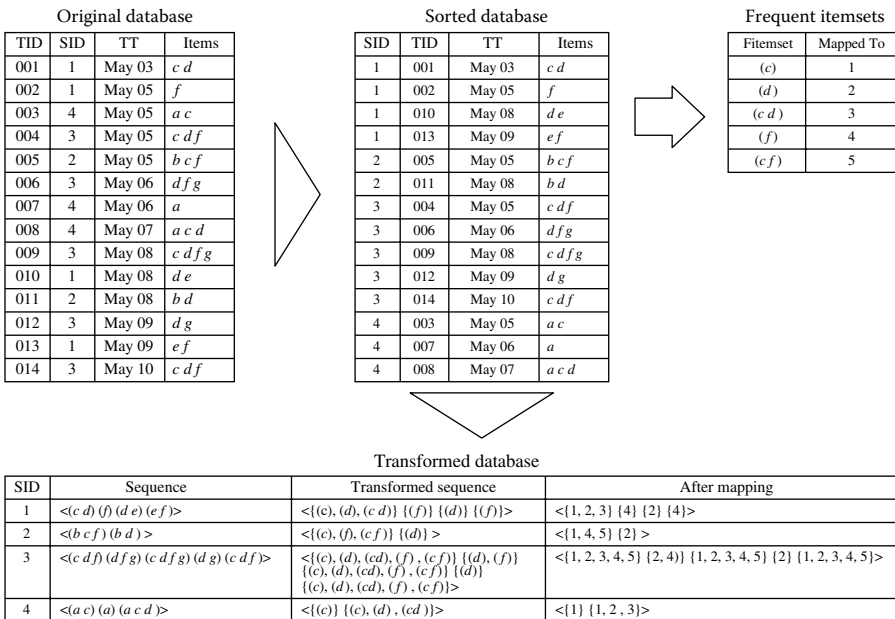
**Figure 4.3** Transformation from the original database to the transformed database.

TABLE 4.4 Frequent Sequences and Candidate Sequences

F_1	C_2	F_2	C_3	F_3	C_4	F_4
(1)	(11) (12) (21) (13) (31)	(11) (12)	(111) (112) (113) (114)	(124) (142)	(1244)	(1424)
(2)	(14) (41) (15) (51) (22)	(13) (14)	(122) (132) (124) (142)	(144) (224)	(1424)	(2424)
(3)	(23) (32) (24) (42) (25)	(22) (32)	(134) (144) (222) (224)	(242) (324)	(2244)	(3424)
(4)	(52) (33) (34) (43) (35)	(24) (42)	(242) (322) (324) (342)	(342) (244)	(2424)	
(5)	(53) (44) (45) (54) (55)	(52) (34)	(244) (422) (424) (442)	(424) (344)	(3244)	
		(44)	(522) (344) (444)		(3424)	

transformation phase on this example. In the sort phase, transactions in the database are sorted with sequence-id (SID) as the major key and transaction-time (TT) as the minor key. Then, in the itemset phase, itemsets are derived in the similar manner to Apriori. Note that the support of an itemset is the number of transaction sequences, including the itemset, but not the number of transactions including it. Thus, the resulting set of frequent 1-itemsets in this case is $\{c, d, f\}$. In the transformation phase, each transaction sequence is transformed into a list of sets of itemsets as shown in the bottom of Figure 4.3 by replacing each transaction in the sequence with a set of itemsets the transaction contains. Note that the second transaction is dropped in the transaction sequence 4 because it consists of only one nonfrequent itemset $\{a\}$.

AprioriAll generates a set of candidate sequences C_2 from F_1 by calling the function *apriori-gen-2*. The resulting C_2 is shown in Table 4.4. The function *apriori-gen-2* is similar to *apriori-gen*, but differs in its join operation: The join operation of *apriori-gen-2* generates two new k -sequences from two $(k - 1)$ -sequences whenever they are joinable, while the join operation of *apriori-gen* generates only one k -itemset from two $(k - 1)$ -itemsets. For example, when deriving C_2 , both two sequences (12) and (21) are generated from (1) and (2). In addition, (11) is also generated by joining the identical sequence (1). This is necessary to generate a sequence in which multiple occurrences of an itemset is allowed.

Counting the support of each candidate sequence is done in the similar way as Apriori using a hash-tree, and F_2 , a set of frequent 2-sequences, is derived as shown in Table 4.4. This F_2 is used to generate a set of candidate sequences C_3 as well. Note that from (11) and (12), a 3-sequence (112) is generated by joining them, but not (121) because its subsequence (21) is not included in F_2 . This process consisting of the candidate generation and support counting is repeated until no more frequent sequences are derived. In this example, since no candidate of 5-sequences can be generated from F_4 , F_5 becomes empty and thus, the iteration terminates. Finally, AprioriAll outputs (1424), (2424), (3424), (11), (13), and (52) as the maximal frequent sequences as the other frequent sequences are included in one of them.

4.4.2 Performance Evaluation

In this section, we discuss the performance of Apriori with respect to its runtime, the number of derived association rules and frequent itemsets when *minsup*, *minconf*, and the number of transactions are varied. We used the implementation by Christian Borgelt [1] for this assessment because it provides options that allow us to simulate a




Figure 4.4 Runtime for various *minsup* and *minconf* values.

naive implementation closest to the original Apriori. Thus, we disabled its functions of sorting items with respect to their support and of filtering unused items from transactions.

As a benchmark dataset, we used the Mushroom dataset downloadable from UCI Machine Learning Repository [8], which contains 8124 cases with 23 nominal attributes including a class attribute. Each case is regarded as a transaction, and each attribute value of each case is converted into an item by joining it with the corresponding attribute name, for example, “cap-shape=x,” where cap-shape is an attribute name and x is an attribute value. In 2480 cases, the attribute value of one attribute is missing. Since we ignored missing values, the transactions corresponding to them have 22 items, while the others have 23 items. Some attribute values have different meanings for different attributes. For example, “n” means “none” for the attribute “odor,” while “brown” for “cap-color.” As a result, the number of valid pairs of attribute name and attribute value, that is, number of distinct items, became 118.

First, we show the runtime of Apriori for various *minsup* and *minconf* values in Figure 4.4. All runtimes shown in this section were measured on a PC running Windows XP with 2.8 GHz Pentium IV and 4 GB memory. In these experiments, the maximal number of items per rule is set to 5 for convenience. We also limited the minimal number of items per rule to 2 in order to prevent a rule with no premise from being derived. In addition, a prefix tree was not used to store transactions. From the results, it is obvious that the change of *minconf* does not affect the runtime so much, but the runtime exponentially increases as *minsup* becomes smaller. The similar tendency is observed in Figure 4.5, showing the relation between *minsup* and the number of derived association rules. This is because the number of frequent itemsets exponentially increases as *minsup* becomes smaller, as shown in Figure 4.6.




Figure 4.5 Number of association rules derived for various *minsup* and *minconf* values.

These results show that *minsup*, or the antimonotonicity property of itemsets, is very effective to prune nonfrequent itemsets.

Next, we show the relation between the runtime and the number of transactions in Figure 4.7. In this evaluation, we copied the original dataset multiple times (up to




Figure 4.6 Number of frequent itemsets for various *minsup* values.




Figure 4.7 Runtime for various sizes of the dataset ($minsup = 5$).

4 times). Note that the fraction of each item remains the same for all datasets, so is the number of resulting association rules (frequent itemsets). Figure 4.7 shows that the runtime linearly increases as the number of transactions becomes larger. Consequently, under a certain distribution of items, $minsup$ is much more influential to the runtime than both $minconf$ and the number of transactions in Apriori.

Finally, we briefly mention association rules mined through the experiments, especially, for convenience, those which have only one item representing the class attribute in the consequent. The class value is either “edible” (e) or “poisonous” (p). A typical rule found under $minsup = 0.3$ and $minconf = 0.9$ is “odor = n gill-size = b ring-number = o \Rightarrow class = e,” which is the simplest one among those whose consequent is “class = e,” confidence is 1.0, and support is maximum (0.331). This rule means a mushroom is edible if its odor is none, the size of its gill is broad, and the number of its rings is one. The attributes “odor” and “gill-size” appear as the first and the third test nodes, respectively, in the decision tree learned from this dataset by J48, a decision tree learner available in Weka, under its default setting. A similar rule “odor = n spore-print-color = w gill-size = b \Rightarrow class = e” can be derived from the decision tree and its confidence is 1.0, too, but it is true for only 528 cases, while the association rule is true for 2689 cases. On the other hand, no rule whose confidence is 1.0 and consequent is “class = p” was found under this setting because $minsup$ was too high. When setting $minsup = 0.2$, 470 such rules were found.

In general we can obtain a small number of association rules in a short runtime for a high $minsup$, but many of them could be trivial. To find more interesting rules, we have to use a smaller $minsup$, but it leads to an unacceptable runtime and a huge number of association rules, which in turn would make it harder to find interesting association

rules. More efficient algorithms and better measures are required to find frequent itemsets and interesting association rules, which are the topics of the next section.

4.5 Advanced Topics

Since the first proposal of frequent pattern and association rule mining algorithm by Agrawal and Srikant, there have been many publications on various kinds of improvements, extensions, and applications, ranging from efficient scalable data mining methodologies, to handling a wide diversity of data types, various extended mining tasks, and a variety of new applications. Some of the important advanced topics are briefly described in this section. There are good tutorials and surveys for frequent pattern mining by Han et al. [16] and Goethals [15] that contain a substantial amount of references.

4.5.1 Improvement in Apriori-Type Frequent Pattern Mining

There have been many attempts to devise more efficient algorithms of frequent itemset mining in the framework of Apriori algorithm in that they generate candidates. These include hash-based technique, partitioning, sampling, and using vertical data format.

Hash-based technique can reduce the size of candidate itemsets. Each itemset is hashed into a corresponding bucket by using an appropriate hash function. Since a bucket can contain different itemsets, if its count is less than a minimum support, these itemsets in the bucket can be removed from the candidate sets. DHP [26] uses this idea.

Partitioning can be used to divide the entire mining problem into n smaller ones [29]. The dataset is divided into n nonoverlapping partitions such that each partition fits into main memory and each partition is mined separately. Since any itemset that is potentially frequent must occur as a frequent itemset in at least one of the partitions, all the frequent itemsets found this way are candidates, which can be checked by accessing the entire dataset only once.

Sampling is simply to mine a random sampled small subset of the entire data. Since there is no guarantee that we can find all the frequent itemsets, normal practice is to use a lower support threshold. Trade-off has to be made between accuracy and efficiency.

Vertical data format associates TID with each itemset, whereas Apriori uses a horizontal data format, that is, frequent itemsets are associated with each transaction. With the vertical data format, mining can be performed by taking the intersection of TIDs. The support count is simply the length of the TID set for the itemset. There is no need to scan the database because TID set carries the complete information required for computing support. This technique requires,

Algorithm 4.5 FP-Growth Algorithm: $F[I](\text{FP-tree})$

```

 $F[I] = \emptyset;$ 
foreach  $i \in \mathcal{I}$  that is in  $\mathcal{D}$  in frequency increasing order do begin
     $F[I] = F[I] \cup \{I \cup \{i\}\};$ 
     $\mathcal{D}^i = \emptyset;$ 
     $H = \emptyset;$ 
    foreach  $j \in \mathcal{I}$  in  $\mathcal{D}$  such that  $j < i$  do begin
        // ( $j$  is more frequent than  $i$ )
        Select  $j$  for which support  $(I \cup \{i, j\}) \geq \text{minsup};$ 
         $H = H \cup \{j\};$ 
    end
    foreach  $(Tid, X) \in \mathcal{D}$  with  $i \in X$  do
         $\mathcal{D}^i = \mathcal{D}^i \cup \{(Tid, \{X \setminus \{i\}\} \cap H)\};$ 
        Construct conditional FP-tree from  $\mathcal{D}^i;$ 
        Call  $F[I \cup \{i\}](\text{conditional FP-tree});$ 
         $F[I] = F[I] \cup F[I \cup \{i\}](\text{conditional FP-tree});$ 
    end

```

given a set of candidate itemsets, that their TIDs are available in main memory, which is of course not always the case. However, it is possible to significantly reduce the total size by using a depth-first search. Eclat [43] uses this strategy. In the depth-first approach, it is necessary to store at most the *TID* list of all k -itemsets with the same first $k - 1$ items ($k - 1$ prefix) at depth d with $k \leq d$ in the main memory.

4.5.2 Frequent Pattern Mining Without Candidate Generation

The most outstanding improvement over Apriori would be a method called *FP-growth* (frequent pattern growth) that succeeded in eliminating candidate generation [17, 18]. It adopts a divide and conquer strategy by (1) compressing the database representing frequent items into a structure called FP-tree (frequent pattern tree) that retains all the essential information and (2) dividing the compressed database into a set of conditional databases, each associated with one frequent itemset and mining each one separately. It scans the database only twice. In the first scan, all the frequent items and their support counts (frequencies) are derived and they are sorted in the order of descending support count in each transaction. In the second scan, items in each transaction are merged into an FP-tree and items (nodes) that appear in common in different transactions are counted. Each node is associated with an item and its count. Nodes with the same label are linked by a pointer called a node-link. Since items are sorted in the descending order of frequency, nodes closer to the root of the FP-tree are shared by more transactions, thus resulting in a very compact representation that stores all the necessary information. Pattern growth algorithm works on FP-tree

by choosing an item in the order of increasing frequency and extracting frequent itemsets that contain the chosen item by recursively calling itself on the conditional FP-tree, that is, FP-tree conditioned to the chosen item. FP-growth is an order of magnitude faster than the original Apriori algorithm. The algorithm of FP-growth is given in Algorithm 4.5. $F[\emptyset](\text{FP-tree})$ returns all the frequent itemsets. As noted easily, the divide and conquer strategy mentioned by Han et al. is equivalent to the depth-first search without candidate generation. The \mathcal{D}^i is called i -projected database and generally much smaller than the FP-tree of the whole database. It is, thus, expected that \mathcal{D}^i fits in the main memory even if the latter does not. The idea of pattern growth can also be applicable to closed itemset mining [27] (see Section 4.5.4) and sequential pattern mining [28] (see Section 4.5.8).

4.5.3 Incremental Approach

When the database is not stationary and a new batch of transactions keeps being added, it happens that some items that were frequent become no more frequent (losers) and some other items that were infrequent become frequent (winners). Rerunning Apriori or any other frequent pattern mining algorithm each time the database is updated is not efficient. The FUP algorithm in [12] provides a way to incrementally update the frequent itemsets using Apriori framework. It works efficiently on the updated database since the size of the increment database $\Delta\mathcal{D}$ is generally much smaller than the initial database \mathcal{D} .

Let F_k, F'_k be the frequent k -itemsets in \mathcal{D} and $\mathcal{D} \cup \Delta\mathcal{D}$, respectively, and C_k be the candidate frequent itemsets in $\mathcal{D} \cup \Delta\mathcal{D}$. At k -th iteration, C_k can be generated from F'_{k-1} using apriori-gen function. Any itemset in F_k that contains any one of the losers of size $k-1$ (those which are in F_{k-1} but not in F'_{k-1}) as its subset are filtered out from F_k without checking $\Delta\mathcal{D}$. Frequency of the remaining itemsets in F_k are counted over $\Delta\mathcal{D}$ and those frequent in $\mathcal{D} \cup \Delta\mathcal{D}$ are identified (A), and excluded from C_k because we know that they are frequent. The remaining itemsets are those not in F_k . Their frequency is counted over $\Delta\mathcal{D}$ and those not frequent in $\Delta\mathcal{D}$ are removed from C_k because we know that they are infrequent in \mathcal{D} . Frequency of the remaining elements in C_k are counted over $\mathcal{D} \cup \Delta\mathcal{D}$ and the frequent ones are retained (B). F'_k is $A \cup B$. As can be seen above, FUP has to scan the updated database for each k , but the size of the C_k is expected to be very small. The experiment shows that it is only about 2 to 5% of that of rerunning Apriori for the updated database, and FUP runs 2 to 16 times faster than Apriori.

4.5.4 Condensed Representation: Closed Patterns and Maximal Patterns

An itemset (pattern) X is a maximal itemset if (1) there exists no itemset X' such that X' is a proper superset of X . An itemset (pattern) X is a closed itemset if (1) there exists no itemset X' such that X' is a proper superset of X and (2) every transaction containing X also contains X' . They are frequent if their support is no less than the *minsup*. A closed itemset satisfies $I(\mathcal{T}(X)) = X$, where $\mathcal{T}(X) = \{t \in \mathcal{D} | X \subseteq t\}$ and $I(S) = \bigcap_{t \in S} t$ for $S \subseteq \mathcal{D}$. For any two itemsets X and Y , if $X \subset Y$ and their support

is the same, X is not a closed itemset. A closed itemset is a lossless representation, whereas a maximal itemset is not. Thus, once the closed itemsets are found, all the frequent itemsets can be derived from them. A rule $X \Rightarrow Y$ is an association rule on frequent closed itemsets if (1) both X and $X \cup Y$ are frequent closed itemsets, (2) there does not exist a frequent closed itemset Z such that $X \subset Z \subset (X \cup Y)$, and (3) the confidence of the rule is no less than *minconf*. The complete set of association rules can be generated once frequent closed itemsets are found.

CLOSET partitions the database and decomposes the problem into a set of subproblems, each with the corresponding conditional database, and it is known efficient [27]. First, all the frequent items are derived and sorted in the order of descending support count as $f_list = \langle i_1, i_2, \dots, i_n \rangle$. The j -th subproblem ($1 \leq j \leq n$) is to find the complete set of frequent closed itemsets containing i_{n+1-j} but no i_k (for $n+1-j < k \leq n$). The i_{n+1-j} conditional database is the subset of transactions containing i_{n+1-j} , where all the occurrences of infrequent items, item i_{n+1-j} , and items following i_{n+1-j} in the f_list are omitted. The corresponding FP-tree is generated and used for search. Each subproblem is recursively decomposed if necessary. The frequent closed itemsets are identified from the conditional database using the following properties. If X is a frequent closed itemset, there is no item appearing in every transaction in the X -conditional database. If an itemset Y is the maximal set of items appearing in every transaction in the X -conditional database, and $X \cup Y$ is not subsumed by some already found frequent closed itemset with identical support, $X \cup Y$ is a frequent closed itemset. As in FP-growth, further optimization is possible.

LCM is another algorithm, known to be the most efficient, to find the closed patterns (itemsets) [34]. It derives frequent closed itemsets via a closure operation without generating nonclosed itemsets. A closure of an itemset X , denoted by $Clo(X)$, is the unique smallest closed itemset including X , that is, $I(\mathcal{T}(X))$. Without loss of generality, we assume all items in a transaction database are uniquely indexed by contiguous natural numbers. Then, $X(i) = X \cap \{1, \dots, i\}$ is called the i -prefix of X , which is the subset of X having only elements no greater than i . The *core index* of a closed itemset X , denoted by $core_i(X)$, is the minimum index i such that $\mathcal{T}(X(i)) = \mathcal{T}(X)$. LCM generates, from a frequent closed itemset X , another frequent closed itemset Y such that $Y = Clo(X \cup \{i\})$ and $X(i-1) = Y(i-1)$, where i is an item that satisfies $i \notin X$ and $i > core_i(X)$. Y is called the *prefix-preserving closure extension*, or ppc-extension for short, of X . LCM recursively applies this closure operation to closed itemsets from an empty itemset to larger ones in a depth-first manner. Completeness and nonredundancy of the enumeration of closed itemsets by LCM are guaranteed by the following property: If Y is a nonempty closed itemset, then there is just one closed itemset X such that Y is a ppc-extension of X . Since LCM generates a new frequent closed itemset Y from $\mathcal{T}(X)$ and a subset of \mathcal{I} , its time complexity to enumerate all frequent closed itemsets for X is $O(|\mathcal{T}(X)| \times |\mathcal{I}|)$, where $|\mathcal{T}(X)|$ is the summation of size of each transaction included in $\mathcal{T}(X)$. Let \mathcal{C} be a set of all frequent closed itemsets in \mathcal{D} . Then, the time complexity of LCM is linear in $|\mathcal{C}|$ with a factor depending on $|\mathcal{T}| \times |\mathcal{I}|$. In fact, to improve the computation time and memory use, LCM incorporates three techniques: occurrence deliver, anytime database reduction, and fast prefix-preserving test. Occurrence deliver constructs

$\mathcal{T}(X \cup \{i\})$ for all i by scanning $\mathcal{T}(X)$ only once instead of scanning it for each i . Anytime database reduction reduces the size of the database by removing unnecessary transactions and items from it each time before an iteration starts with the current closed itemset to reduce both the computation time and memory use. Fast prefix-preserving test significantly reduces the number of items to be accessed to test the equality $X(i-1) = Y(i-1)$ by checking only items j such that $j < i$, $j \notin X(i-1)$ and they are included in the transaction of the minimum size in $\mathcal{T}(X \cup \{i\})$ instead of actually generating a closure when performing a ppc-extension. If an item j is included in every transaction in $\mathcal{T}(X \cup \{i\})$, then j is included in $Clo(X \cup \{i\})$, thus $X(i-1) \neq Y(i-1)$.

4.5.5 Quantitative Association Rules

When the item has a continuous numeric value, current frequent itemset mining algorithms are not applicable unless the values are discretized and appropriate intervals defined. This is known as *quantitative frequent itemset (QFI) mining*. The items can be both categorical and numeric. An example is $\{\langle \text{Age: } [30,39] \rangle, \langle \text{House-owner: Yes} \rangle, \langle \text{Married: Yes} \rangle\}$, where an item is represented as $\langle \text{attribute: its value (range)} \rangle$. QFI mining was initially proposed in the study of mining quantitative association rules [31], but later density-based subspace clustering has commonly been applied because a QFI is viewed as an axis-parallel hyper-rectangular containing a cluster of transactions in a numeric attribute space. SUBCLUE [20] and QFIMiner [36] are two such examples. QFIMiner finds in $O(N \log N)$ all dense clusters of no less than *minsup* in all subspaces formed by both numeric and categorical attributes, where N is the number of transactions. An optimal value interval for each numeric item in each frequent itemset is obtained by Apriori-like level-wise algorithm with the antimonicity property of dense clusters. QFIMiner is shown to be faster than SUBCLUE and scales very well.

4.5.6 Using Other Measure of Importance/Interestingness

The problem of support-confidence framework is that there is no valid means to determine appropriate values for *minsup* and *minconf*. Especially setting *minsup* too high will miss important rules and setting it too low will generate too many rules. In fact, it is possible that a rule with infrequent itemsets is of great interest for some applications. Further, this framework fails to capture the notion of correlation. It can happen that a rule $X \Rightarrow Y$ which satisfies both *minsup* and *minconf* constraints has no correlation between X and Y , that is, $\text{support}(X) \times \text{support}(Y) = \text{support}(X \cup Y)$.

Therefore, an alternative approach is to use other measures that account for importance or interestingness of a rule and select rules that have high score for these measures. Support and confidence can still be used as a constraint (setting *minsup* and *minconf* to 0 means not to use them at all). These measures include lift, leverage, redundancy, productivity, and well-known statistical measures such as chi-square, correlation coefficient, information gain, and so on.

Lift and leverage represent the ratio and the difference between the support and the support that would be expected if X and Y were independent, respectively. They try

to find rules with strong correlations between X and Y .

$$\begin{aligned}\text{lift}(X \Rightarrow Y) &= \frac{\text{confidence}(X \Rightarrow Y)}{\text{confidence}(\emptyset \Rightarrow Y)} = \frac{\text{support}(X \Rightarrow Y)}{\text{support}(X) \times \text{support}(Y)} \\ \text{leverage}(X \Rightarrow Y) &= \text{support}(X \Rightarrow Y) - \text{support}(X) \times \text{support}(Y) \\ &= \text{support}(X) \times (\text{confidence}(X \Rightarrow Y) - \text{support}(Y))\end{aligned}$$

Redundant rule constraint discards a rule $X \Rightarrow Y$ if $\exists Z \in X : \text{support}(X \Rightarrow Y) = \text{support}(X - [Z] \Rightarrow Y)$. A more powerful constraint is productive constraint. A rule is said to be productive if its improvement is greater than 0, where the rule's improvement is defined as

$$\text{improvement}(X \Rightarrow Y) = \text{confidence}(X \Rightarrow Y) - \max_{Z \subset X} (\text{confidence}(Z \Rightarrow Y)).$$

The improvement of a redundant rule cannot be greater than 0 and hence a constraint that rules must be productive discards all redundant rules. Further, it can discard rules that include items in the antecedent that are independent of the consequent, given the remaining items in the antecedent.

Statistical measures are useful in finding discriminative patterns (itemsets). However, these measures do not satisfy the antimonotonicity property, and finding the best k patterns or rules is not that easy. If a measure is convex with respect to its arguments, it is possible to estimate its upperbound for supersets of a pattern X (itemset) for a fixed conclusion Y (normally, a class value) [23] and use this to prune the search space. Statistical measures mentioned above satisfy this property.

Webb's KORD algorithm [39] finds k -optimal rules through the space of pairs X and Y (without fixing Y) and uses leverage as a measure to optimize using various pruning strategies.

4.5.7 Class Association Rules

When a transaction t is associated with a class cl , it is natural to use association rules for classification purpose. The association rules mined for classification purpose are called *class association rules* (CARs). CARs have the form $\{\langle p_1 : q_1 \rangle, \langle p_2 : q_2 \rangle, \dots, \langle p_m : q_m \rangle\} \Rightarrow cl$. Here a numeric item has a numeric interval value, whereas a categorical item has a categorical value. Let \mathcal{D}_{cl} be a set of all instances having a class cl in \mathcal{D} . CBA [22], CMAR [21], and CAEP [14] are the representative CAR-based classification systems. Especially, CAEP introduces a notion of emergent patterns and uses the strength of all CARs. Let the support of an itemset a by \mathcal{D}_{cl} be $\text{support}_{\mathcal{D}_{cl}}(a) = |\{t \in \mathcal{D}_{cl} | a \in t\}| / |\mathcal{D}_{cl}|$. A set of QFIs, $\text{FQFI}(cl)$, in which every itemset a satisfies $\text{support}_{\mathcal{D}_{cl}}(a) \geq \text{minsup}$, is derived for every cl from \mathcal{D}_{cl} . Next, for every $a \in \text{FQFI}(cl)$, the *growth rate* defined by $\text{growth_rate}_{\overline{\mathcal{D}}_{cl} \rightarrow \mathcal{D}_{cl}}(a) = \text{support}_{\mathcal{D}_{cl}}(a) / \text{support}_{\overline{\mathcal{D}}_{cl}}(a)$ is calculated for each class cl , where $\overline{\mathcal{D}}_{cl} = \mathcal{D} - \mathcal{D}_{cl}$ represents the opponent instances of cl . When the growth rate of a is not less than its threshold $\rho (\geq 1)$, that is, $\text{growth_rate}_{\overline{\mathcal{D}}_{cl} \rightarrow \mathcal{D}_{cl}}(a) \geq \rho$, a is called an *emergent pattern* (EP) and is selected for a rule body where its head is the class cl , that is, $a \Rightarrow cl$. Let $\text{FEP}(cl)$ be a set

of all EPs selected from $FQFI(cl)$ under this measure. The underlying principle here is to select the rule bodies that are strong enough to differentiate the class cl from the others. The strength of an EP a is measured by the relative difference between $\text{support}_{\mathcal{D}_{cl}}(a)$ and $\text{support}_{\overline{\mathcal{D}}_{cl}}(a)$: $\text{support}_{\mathcal{D}_{cl}}(a)/(\text{support}_{\mathcal{D}_{cl}}(a) + \text{support}_{\overline{\mathcal{D}}_{cl}}(a)) = \text{growth_rate}_{\overline{\mathcal{D}}_{cl} \rightarrow \mathcal{D}_{cl}}(a)/(\text{growth_rate}_{\overline{\mathcal{D}}_{cl} \rightarrow \mathcal{D}_{cl}}(a) + 1)$. This can be aggregated to define the *aggregate score* defined by $\text{score}(t, cl) = \sum_{a \subseteq t, a \in FEP(cl)} \frac{\text{growth_rate}(a)}{\text{growth_rate}(a) + 1} * \text{support}_{\mathcal{D}_{cl}}(a)$ which represents the possibility of t to be classified into cl by EPs in $FEP(cl)$. Since the distribution of the number of EPs is not uniform over cl , instances may get higher scores for some classes. Another factor, called a *base score*, which is defined to be the median of all aggregate scores in $\{\text{score}(t, cl) | t \in \mathcal{D}_{cl}\}$, is introduced to offset this bias, giving the normalized score defined by $\text{norm_score}(t, cl) = \frac{\text{score}(t, cl)}{\text{base_score}(cl)}$. The cl for which the normalized score is maximum is assigned to the class of t . This was shown to perform very well.

The problem with CAEP is that it discretizes each numeric attribute by an entropy measure without taking account of the dependency that exists in multiple attributes, and thus a cluster of instances having the same class can often be fragmented. Natural solution is to combine QFIMiner and CAEP, which is LSC-CAEP [37, 36].

4.5.8 Using Richer Expression: Sequences, Trees, and Graphs

Mining frequent itemsets started with a simple transaction dataset, but later it has been generalized to be able to deal with richer expression such as sequences, trees, and graphs. The pioneering work to mine sequential patterns by Agrawal and Srikant has already been discussed in Section 4.2.2. PrefixSpan [28] is another representative algorithm in frequent sequential pattern mining, which is a pattern-growth based algorithm and adopts a divide and conquer strategy similar to FP-growth to avoid unfruitful enumeration of smaller candidates to find larger patterns. PrefixSpan, first, finds sequential patterns consisting of only one item, and then, for each of them, say i_k , extracts a set of sequences containing it, that is, the $\langle i_k \rangle$ -projected database. From each such projected database, PrefixSpan finds frequent sequential patterns of size 2 having $\langle i_k \rangle$ as their prefix, and again generates a projected database for each size 2 pattern newly found to find sequential patterns of size 3. This process is recursively repeated until no more sequential patterns are found.

A tree is characterized by V , a set of vertices, and E , a set of edges. A labeled tree assigns a set of labels L to either one or both of vertices and edges. An edge connects a vertex to another one. Every two vertices in a tree are reachable through one or more edges, but there is no cyclic path. TreeMinerV [44] and FREQT [7] are representative algorithms to mine subtrees frequently appearing in a collection of trees. They were independently proposed, but share the same level-wise strategy to enumerate frequent subtrees, which finds frequent subtrees having $k + 1$ vertices ($(k + 1)$ -subtrees) from k -subtrees by adding one edge to every possible position on a specific path called the rightmost path of each k -subtree with a vertex corresponding to the other end of the edge. Dryade [33] is a tree mining algorithm that can find frequent closed subtrees. A closed subtree is a maximal subtree among those having the same frequency. Unlike

the other tree mining algorithms, Dryade assembles frequent closed subtrees level by level from a set of basic units called *tiles*, which are one depth closed subtrees.

A *graph* is a super class of trees and can have cyclic paths. AGM [19] is the first algorithm that mines frequent subgraphs from a collection of graphs by a complete search. It is based on Apriori and generates a candidate subgraph of size k (k -subgraph) from two known frequent $(k - 1)$ -subgraphs which share the same $(k - 2)$ -subgraph. Since there is no edge information available between the two $(k - 1)$ -th vertices, all the possibilities are considered. AGM generates two k -subgraphs from a pair of $(k - 1)$ -subgraphs, one with an edge between them and the other without an edge (this is a case where there are no labels defined for edges). Although Apriori-based approach enables to conduct a systematic complete search of frequent subgraphs, it has to generate a large number of candidates that do not actually exist in a given set of graphs. AGM uses adjacency matrix to represent a graph and introduces a notion of canonical form to solve subgraph isomorphism which is known to be NP-complete. gSpan [41] is one representative pattern-growth-based subgraph mining algorithm. It finds frequent subgraphs in a depth-first manner by adding an edge to each possible position on the rightmost path of a known frequent subgraph. gSpan takes into account only the edges that actually exist in a given set of graphs, so it never generates candidates that do not actually exist. GBI [42] and SUBDUE [13] are greedy algorithms to find frequent subgraphs, which recursively replace every occurrence of a typical subgraph in a graph with a new vertex. The typicality is defined by a measure based on frequency, for example, information gain in GBI and the minimum description length in SUBDUE. DT-CIGBI [25] generates a decision tree that classifies unknown graphs from a set of training graphs with known classes. It invokes a graph mining algorithm, CI-GBI [24], an extension of GBI, at every test node of the decision tree. The resulting frequent subgraphs are used as attributes of graphs, and the most discriminative one is chosen to split the set of graphs that reached the node into two subsets: those which include the subgraph and the others.

4.6 Summary

Experimenting with Apriori-like algorithm is the first thing that data miners try to do. In this chapter the basic concepts and algorithms of Apriori family (Apriori, AprioriTid, AprioriAll) were introduced first and then their working mechanisms were explained with illustrative examples, followed by a performance evaluation of Apriori using a typical freely available implementation. Since Apriori is so fundamental and easy to implement, there are many variants of it. The limitation of Apriori approach is discussed and an overview of recent important advancement in frequent pattern mining methodologies is provided. There are other topics that cannot be covered in this chapter. These include use of constraints, colossal patterns, noise handling, and top-k representatives.

4.7 Exercises

1. Prove that Apriori can derive all frequent itemsets from a given transaction database.
2. Prove the following relation:

$$\text{support}(X \cup Y \cup Z) \geq \text{support}(X \cup Y) + \text{support}(X \cup Z) - \text{support}(X),$$

where X , Y , and Z are itemsets in a database.

3. Given the database shown in Table 4.5, find all frequent itemsets using Apriori and AprioriTid for $\text{minsup} = 0.3$ and compare their efficiency.
4. Explain the relation between a hash-tree and a trie.
5. Draw an FP-tree for the database shown in Table 4.5 and explain how frequent itemsets are derived from the FP-tree.
6. Download and install Weka on your computer, and mine association rules by using Apriori from the Soybean dataset included in the Weka's package for various metrics to evaluate association rules using the same minimum threshold (fix the other parameters). Then, report how the resulting association rules change according to the metrics.
7. Draw a prefix tree to store the database in Section 4.4 with reference to [10] and explain how the efficiency of frequency counting can be improved in this case.
8. In an FP-tree, items in a transaction are sorted in the order of descending support count, while in a prefix tree for Apriori they are sorted in the order of ascending support count. Discuss the reason why they adopt the different orders.
9. When a transaction database has a small number of very long transactions, Apriori-based algorithms take much time to mine frequent itemsets. Explain the reason why they need so much time and propose an efficient method of mining closed itemsets from such a database.

TABLE 4.5 Database for Exercise 3

TID	Items
T01	Cheese, Milk, Egg
T02	Apple, Cheese
T03	Apple, Bread, Cheese, Orange, Grape
T04	Bread, Egg, Orange
T05	Cheese, Milk, Grape
T06	Apple, Cheese, Egg, Orange
T07	Bread, Cheese, Orange
T08	Cheese, Egg, Grape
T09	Bread, Cheese, Egg, Grape
T10	Bread, Cheese, Grape

TABLE 4.6 Sequence Database
for Exercise 10

SID	Transaction Sequences
S01	$\langle (bc)(d)(ab)(def) \rangle$
S02	$\langle (abc)(cf)(df) \rangle$
S03	$\langle (cef)(df)(ab)(f) \rangle$
S04	$\langle (be)(ac)(cdf) \rangle$

10. Given the sequence database shown in Table 4.6, find frequent sequential patterns by AprioriAll for $minsup = 0.5$.

References

- [1] <http://www.borgelt.net/apriori.html>.
- [2] <http://www.cs.bme.hu/~bodon/en/apriori/>.
- [3] R. Agrawal and R. Srikant. Fast algorithms for mining association rules. In *Proc. of the 20th International Conference on Very Large Data Bases (VLDB 1994)*, pages 487–499, 1994.
- [4] R. Agrawal and R. Srikant. Fast algorithms for mining association rules. IBM Research Report RJ9839, IBM Research Division, Almaden Research Center, 1994.
- [5] R. Agrawal and R. Srikant. Mining sequential patterns. IBM Research Report RJ9910, IBM Research Division, Almaden Research Center, 1994.
- [6] R. Agrawal and R. Srikant. Mining sequential patterns. In *Proc. of the 11th International Conference on Data Engineering (ICDE 1995)*, pages 3–14, 1995.
- [7] T. Asai, K. Abe, S. Kawasoe, H. Arimura, H. Sakamoto, and S. Arikawa. Efficient substructure discovery from large semi-structured data. In *Proc. of the 2nd SIAM International Conference on Data Mining*, pages 158–174, 2002.
- [8] C. Blake and C. Merz. UCI repository of machine learning databases, 1998. <http://www.ics.uci.edu/~mllearn/MLRepository.html>.
- [9] F. Bodon. Surprising results of trie-based fim algorithms. In *Proc. of the IEEE ICDM Workshop on Frequent Itemset Mining Implementations (FIMI'04)*, volume 126 of *CEUR Workshop Proceedings*, 2004. <http://ftp.informatik.rwth-aachen.de/Publications/CEUR-WS/Vol-126/bodon.pdf>.
- [10] C. Borgelt. Efficient implementations of Apriori and Eclat. In *Proc. of the IEEE ICDM Workshop on Frequent Itemset Mining Implementations (FIMI'03)*,

- volume 90 of *CEUR Workshop Proceedings*, 2003. <http://ftp.informatik.rwth-aachen.de/Publications/CEUR-WS/Vol-90/borgelt.pdf>.
- [11] S. Brin, R. Motwani, J. D. Ullman, and S. Tsur. Dynamic itemset counting and implication rules for market basket data. In *Proc. of ACM SIGMOD International Conference on Management of Data (SIGMOD 1997)*, pages 255–264, 1997.
 - [12] D. W. Cheung, J. Han, and C. Y. Wong. Maintenance of discovered association rules in large databases: An incremental updating technique. In *Proc. of the 1996 ACM SIGMOD International Conference on Management of Data*, pages 13–23, 1996.
 - [13] D. J. Cook and L. B. Holder. Substructure discovery using minimum description length and background knowledge. *Journal of Artificial Intelligence Research*, Vol. 1, pages 231–255, 1994.
 - [14] G. Dong, X. Zhang, L. Wong, and J. Li. Caep: Classification by aggregating emerging patterns. In *Proc. of the 2nd International Conference on Discovery Science (DS '99), LNAI 1721*, Springer, pages 30–42, 1999.
 - [15] B. Goethals. Survey on frequent pattern mining, 2003. http://www.adrem.ua.ac.be/bibrem/pubs/fpm_survey.pdf
 - [16] J. Han, H. Cheng, D. Xin, and X. Yan. Frequent pattern mining: Current status and future direction. *Data Mining and Knowledge Discovery*, Vol. 15, No. 1, pages 55–86, 2007.
 - [17] J. Han, J. Pei, and Y. Yin. Mining frequent patterns without candidate generation. In *Proc. of the 2000 ACM SIGMOD International Conference on Management of Data*, pages 1–12, 2000.
 - [18] J. Han, J. Pei, Y. Yin, and R. Mao. Mining frequent patterns without candidate generation: A frequent-pattern tree approach. *Data Mining and Knowledge Discovery*, Vol. 8, No. 1, pages 53–87, 2004.
 - [19] A. Inokuchi, T. Washio, and H. Motoda. General framework for mining frequent subgraphs from labeled graphs. *Fundamenta Informaticae*, Vol. 66, No. 1-2, pages 53–82, 2005.
 - [20] K. Kailing, H. Kriegel, and P. Kroger. Density-connected subspace clustering for high-dimensional data. In *Proc. of the 4th SIAM International Conference on Data Mining*, pages 246–257, 2004.
 - [21] W. Li, J. Han, and J. Pei. Cmar: Accurate and efficient classification based on multiple class-association rules. In *Proc. of the 1st IEEE International Conference on Data Mining (ICDM '01)*, pages 369–376, 2001.
 - [22] B. Liu, W. Hsu, and Y. Ma. Integrating classification and association rule mining. In *Proc. of the 4th International Conference on Knowledge Discovery and Data Mining (KDD-98)*, pages 80–86, 1998.

- [23] S. Morishita and J. Sese. Traversing lattice itemset with statistical metric pruning. In *Proc. of the 19th ACM SIGMOD-SIGACT-SIGART Symposium on Principles of Database Systems (PODS 2000)*, pages 226–236, 2000.
- [24] P. C. Nguyen, K. Ohara, H. Motoda, and T. Washio. CI-GBI: A novel approach for extracting typical patterns from graph-structured data. In *Proc. of the 9th Pacific-Asia Conference on Advances in Knowledge Discovery and Data Mining (PAKDD 2005)*, pages 639–649, 2005.
- [25] K. Ohara, P. C. Nguyen, A. Mogi, H. Motoda, and T. Washio. Constructing decision trees based on chunkingless graph-based induction. In L. B. Holder and D. J. Cook, editors, *Mining Graph Data*, pages 203–226. Wiley-Interscience, 2006.
- [26] J. Park, M. Chen, and P. Yu. An effective hash-based algorithm for mining association rules. In *Proc. of the 1995 ACM SIGMOD International Conference on Management of Data*, pages 175–186, 1995.
- [27] J. Pei, J. Han, and R. Mao. Closet: An efficient algorithm for mining frequent closed itemsets. In *Proc. of the 2000 ACM-SIGMOD International Workshop on Data Mining and Knowledge Discovery*, pages 11–20, 2000.
- [28] J. Pei, J. Han, B. Mortazavi-Asl, H. Pinto, Q. Chen, U. Dayal, and M. C. Hsu. PrefixSpan: Mining sequential patterns efficiently by prefix projected pattern growth. In *Proc. of the 17th International Conference on Data Engineering (ICDE 2001)*, pages 215–224, 2001.
- [29] A. Savasere, E. Omiecinski, and S. Navathe. An efficient algorithm for mining association rules in large databases. In *Proc. of the 21th International Conference on Very Large Data Bases (VLDB 1995)*, pages 432–444. Morgan Kaufmann, 1995.
- [30] R. Srikant and R. Agrawal. Mining generalized association rules. In *Proc. of the 21th International Conference on Very Large Data Bases (VLDB 1995)*, pages 407–419, 1995.
- [31] R. Srikant and R. Agrawal. Mining quantitative association rules in large relational tables. In *Proc. of the 1996 ACM SIGMOD International Conference on Management of Data*, pages 1–12, 1996.
- [32] R. Srikant and R. Agrawal. Mining sequential patterns: Generalization and performance improvement. In *Proc. of the 5th International Conference on Extending Database Technology*, pages 3–17, 1996.
- [33] A. Termier, M. C. Rousset, and M. Sebag. Dryade: A new approach for discovering closed frequent trees in heterogeneous tree databases. In *Proc. of the 4th IEEE International Conference on Data Mining (ICDM '04)*, pages 543–546, 2004.
- [34] T. Uno, T. Asai, Y. Uchida, and H. Arimura. An efficient algorithm for enumerating frequent closed patterns in transaction databases. In *Proc. of the 7th*

- International Conference on Discovery Science (DS '04)*, LNAI 3245, Springer, pages 16–30, 2004.
- [35] T. Washio, H. Matsuura, and H. Motoda. Mining association rules for estimation and prediction. In *Proc. of the 2nd Pacific Asia Conference on Knowledge Discovery and Data Mining (PAKDD 1998)*, pages 417–419, 1998.
 - [36] T. Washio, Y. Mitsunaga, and H. Motoda. Mining quantitative frequent itemsets using adaptive density-based subspace clustering. In *Proc. of the 5th IEEE International Conference on Data Mining (ICDM '05)*, pages 793–796, 2005.
 - [37] T. Washio, K. Nakanishi, and H. Motoda. Deriving class association rules based on levelwise subspace clustering. In *Proc. of the 9th European Conference on Principles and Practice of Knowledge Discovery in Databases (PKDD 2005)*, LNAI 3721, Springer, pages 692–700, 2005.
 - [38] G. Webb. Efficient search for association rules. In *Proc. of the 6th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 99–107, 2000.
 - [39] G. Webb and S. Zhang. K-optimal rule discovery. *Data Mining and Knowledge Discovery*, Vol. 10, No. 1, pages 39–79, 2005.
 - [40] I. H. Witten and E. Frank. *Data Mining: Practical Machine Learning Tools and Techniques, 2nd Edition*. Morgan Kaufmann, San Francisco, 2005. <http://www.cs.waikato.ac.nz/ml/weka/>.
 - [41] X. Yan and J. Han. gSpan: Graph-based substructure pattern mining. In *Proc. of the 2nd IEEE International Conference on Data Mining (ICDM'02)*, pages 721–724, 2002.
 - [42] K. Yoshida and H. Motoda. Clip: Concept learning from inference pattern. *Journal of Artificial Intelligence*, Vol. 75, No. 1, pages 63–92, 1995.
 - [43] M. Zaki. Scalable algorithms for association mining. *IEEE Transactions on Knowledge and Data Engineering*, Vol. 12, No. 3, pages 372–390, 2000.
 - [44] M. J. Zaki. Efficiently mining frequent trees in a forest. In *Proc. of the 8th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD '02)*, pages 71–80, 2002.